

# Introduction to RISC-V ABI

---

PLCT Tech Report

Sinan Lin 林思南

[sinan.lin@aalto.fi](mailto:sinan.lin@aalto.fi)

# Contents of table

- Introduction
- ELF and RISC-V specification
- RISC-V processor-specific ABI

## Step 1. Write down source code

main.c

```
// main.c
#include "add.h"

void _start() {
    asm("li a0, 155\n"
        "li a1, 100\n"
        "call add\n"
        "li a7, 93\n"
        "ecall\n"
        );
}
```

sys\_exit(add(100, 155))

add.h

```
// add.h

int add(int, int);
```

add.c

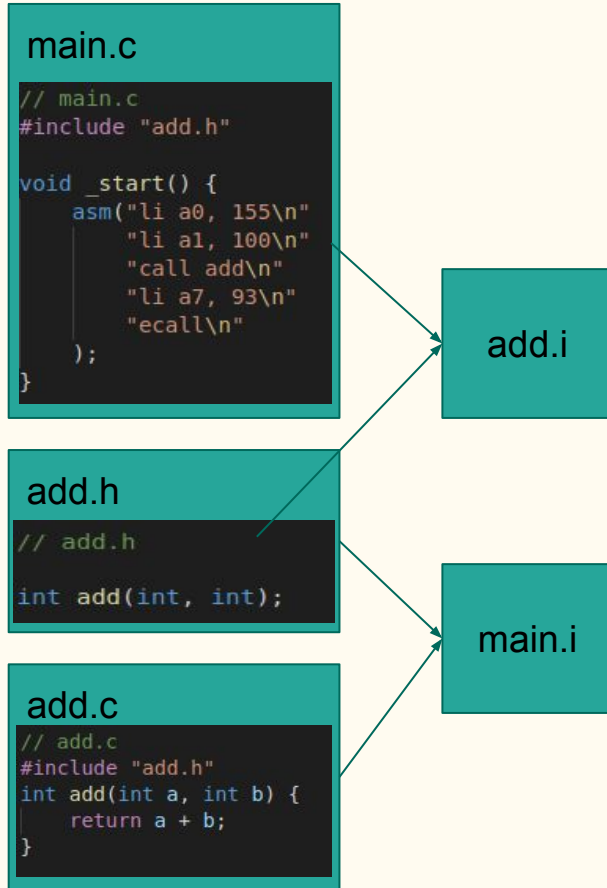
```
// add.c
#include "add.h"
int add(int a, int b) {
    return a + b;
}
```

## Step 2. Use preprocessor to expand macro

Preprocessor expands macros

```
riscv64-unknown-linux-gnu-cpp -P main.c -o main.i
```

```
riscv64-unknown-linux-gnu-cpp -P add.c -o add.i
```



## Step 2. Use preprocessor to expand macro

Preprocessor expands macros

```
riscv64-unknown-linux-gnu-cpp -P main.c -o main.i
```

```
riscv64-unknown-linux-gnu-cpp -P add.c -o add.i
```

### main.c

```
// main.c
#include "add.h"

void _start() {
    asm("li a0, 155\n"
        "li a1, 100\n"
        "call add\n"
        "li a7, 93\n"
        "ecall\n"
    );
}
```

### add.i

```
// add.i
int add(int, int);
int add(int a, int b) {
    return a + b;
}
```

### add.h

```
// add.h
int add(int, int);
```

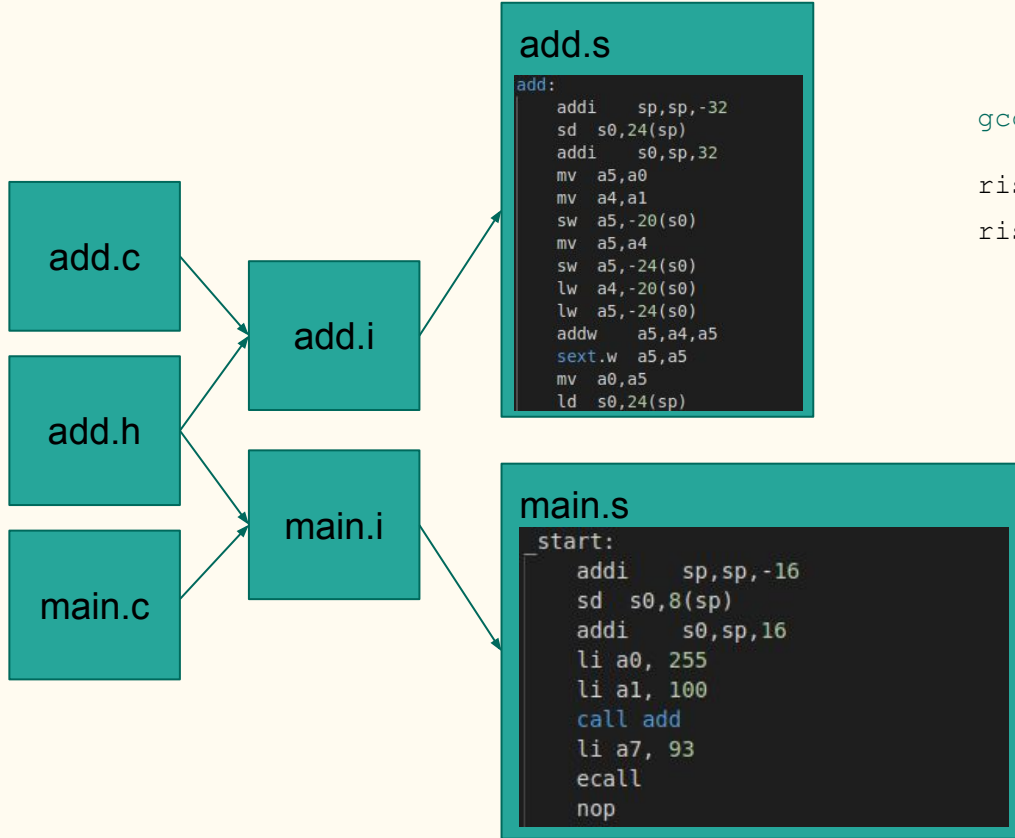
### add.c

```
// add.c
#include "add.h"
int add(int a, int b) {
    return a + b;
}
```

### main.i

```
int add(int, int);
void _start() {
    asm("li a0, 255\n"
        "li a1, 100\n"
        "call add\n"
        "li a7, 93\n"
        "ecall\n"
    );
}
```

## Step 3. GCC compiles source code to assembly



gcc generates assembly


```
riscv64-unknown-linux-gnu-gcc -S main.i -o main.s
riscv64-unknown-linux-gnu-gcc -S add.i -o add.s
```

## Step 4. Assembler converts code to object file

`as` converts assembly code into object code

```
riscv64-unknown-linux-gnu-as main.s -o main.o
```

```
riscv64-unknown-linux-gnu-as add.s -o add.o
```



```
add.o:      file format elf64-littleriscv

Disassembly of section .text:

0000000000000000 <add>:
   0:   fe010113                addi    sp,sp,-32

main.o:      file format elf64-littleriscv


Disassembly of section .text:

0000000000000000 <_start>:
   0:   ff010113                addi    sp,sp,-16
   4:   00813423                sd      s0,8(sp)
   8:   01010413                addi    s0,sp,16
  c:   0ff00513                li      a0,255
 10:   06400593                li      a1,100
 14:   00000097                auipc   ra,0x0
 18:   000080e7                jalr    ra # 14 <_start+0x14>
 1c:   05d00893                li      a7,93
 20:   00000073                ecall
 24:   00000013                nop
 28:   00813403                ld      s0,8(sp)
 2c:   01010113                addi    sp,sp,16
 30:   00008067                ret
```

## Step 5. Linker combines object files and emits executable file

ld converts object files into executable file

```
riscv64-unknown-linux-gnu-ld -static main.o add.o
```



Evaluate result

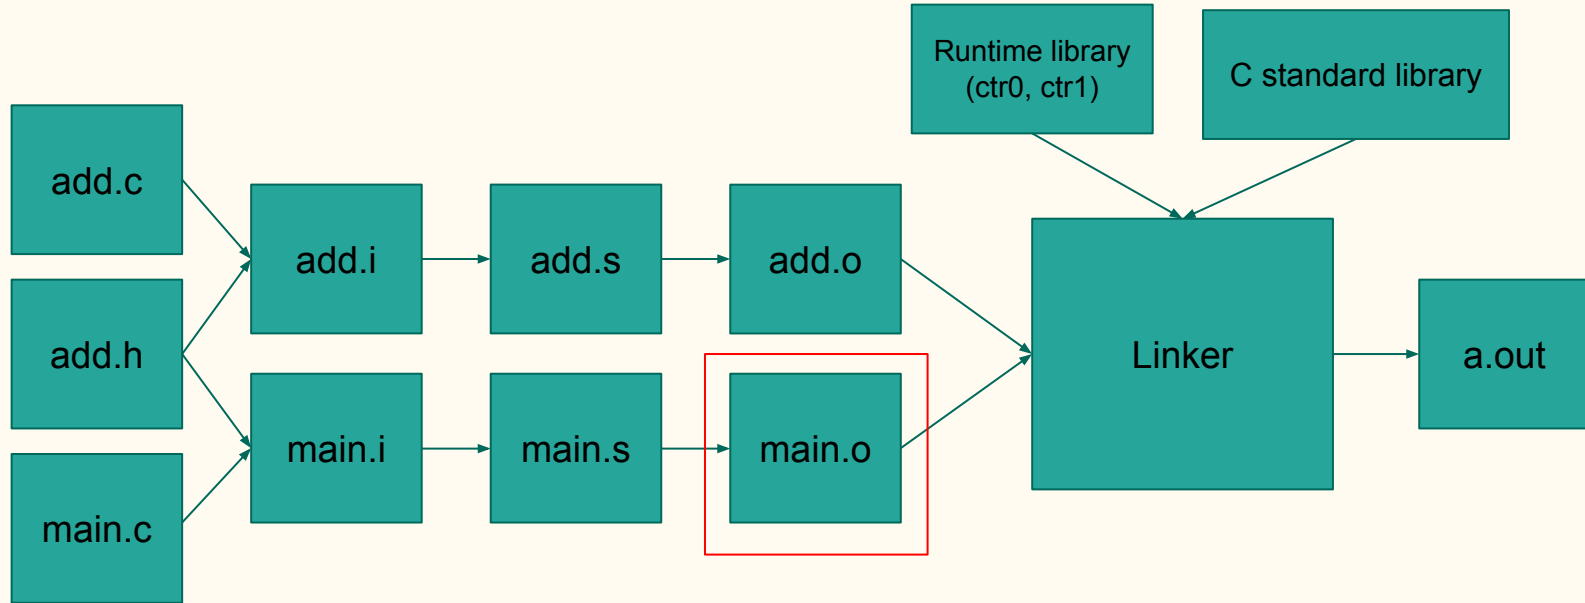
```
qemu-riscv64 a.out
```

```
echo $?
```

```
>> 255
```



- RISC-V specification in ELF format
- RISC-V ABI



# ELF (Executable and linkable file format)

- ELF Header
  - General information about the binary
- Sections
  - Information needed for linking object files in order to build an executable
- Segments
  - Chunks of information to prepare the executable to be loaded into memory.

Sections

ELF Header
.shstrtab
.strtab
.symtab
.comment
.bss
.note
.data
.rela.text
.text
NULL
section header table

Layout of main.o

# ELF (Executable and linkable file format)

- ELF Header
  - General information about the binary
- Sections
  - Information needed for linking object files in order to build an executable
- Segments
  - Chunks of information to prepare the executable to be loaded into memory.

Sections

ELF Header
.shstrtab
.strtab
.symtab
.comment
.bss
.note
.data
.rela.text
.text
NULL
section header table

Layout of main.o

# ELF Header

```
ELF Header:
Magic:  7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
Class:                                ELF64
Data:                                2's complement, little endian
Version:                            1 (current)
OS/ABI:                             UNIX - System V
ABI Version:                         0
Type:                                REL (Relocatable file)
Machine:                             RISC-V
Version:                             0x1
Entry point address:                 0x0
Start of program headers:             0 (bytes into file)
Start of section headers:             496 (bytes into file)
Flags:                                0x5, RVC, double-float ABI
Size of this header:                  64 (bytes)
Size of program headers:              0 (bytes)
Number of program headers:            0
Size of section headers:              64 (bytes)
Number of section headers:            10
Section header string table index:    9
```

main.o

```
typedef struct
{
    unsigned char e_ident[16]; /* Magic number and other info */
    uint16_t e_type; /* Object file type */
    uint16_t e_machine; /* Architecture */
    uint32_t e_version; /* Object file version */
    uint64_t e_entry; /* Entry point virtual address */
    uint64_t e_phoff; /* Program header table file offset */
    uint64_t e_shoff; /* Section header table file offset */
    uint32_t e_flags; /* Processor-specific flags */
    uint16_t e_ehsize; /* ELF header size in bytes */
    uint16_t e_phentsize; /* Program header table entry size */
    uint16_t e_phnum; /* Program header table entry count */
    uint16_t e_shentsize; /* Section header table entry size */
    uint16_t e_shnum; /* Section header table entry count */
    uint16_t e_shstrndx; /* Section header string table index */
} Elf64_Ehdr;
```

riscv-glibc/elf/elf.h

# ELF Header with RISC-V specification

```
ELF Header:
Magic:  7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
Class:                               ELF64
Data:                               2's complement, little endian
Version:                             1 (current)
OS/ABI:                             UNIX - System V
ABI Version:                         0
Type:                                REL (Relocatable file)
Machine:                             RISC-V
Version:                             0x1
Entry point address:                 0x0
Start of program headers:            0 (bytes into file)
Start of section headers:            496 (bytes into file)
Flags:                               0x5, RVC, double-float ABI
Size of this header:                 64 (bytes)
Size of program headers:             0 (bytes)
Number of program headers:           0
Size of section headers:             64 (bytes)
Number of section headers:           10
Section header string table index: 9
```

```
readelf -h main.o
```

RVC allows instructions to be aligned to 16-bit boundaries, the linker is permitted to use RVC instructions such as C.JAL in the relaxation process. ('C' Standard Extension)

ILP64D: Integer calling-convention with hardware floating-point calling convention for FLEN=64

e\_machine

- EM\_RISCV (243) for RISC-V ELF

e\_flags

Bit 0	Bit 1 - 2	Bit 3	Bit 4	Bit 5 - 31
RVC	Float ABI	RVE	TSO	Reserved

- EF\_RISCV\_RVC (0x0001)
- EF\_RISCV\_FLOAT\_ABI\_SOFT (0x0000)
- EF\_RISCV\_FLOAT\_ABI\_SINGLE (0x0002)
- EF\_RISCV\_FLOAT\_ABI\_DOUBLE (0x0004)
- EF\_RISCV\_FLOAT\_ABI\_QUAD (0x0006)
- EF\_RISCV\_FLOAT\_ABI (0x0006)
- EF\_RISCV\_RVE (0x0008)
- EF\_RISCV\_TSO (0x0010)

# Section header table

```
ELF Header:
Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
Class:                                ELF64
Data:                                   2's complement, little endian
Version:                             1 (current)
OS/ABI:                               UNIX - System V
ABI Version:                          0
Type:                                  REL (Relocatable file)
Machine:                              RISC-V
Version:                              0x1
Entry point address:                  0x0
Start of program headers:             0 (bytes into file) 0x1F0
Start of section headers:             496 (bytes into file)
Flags:                                0x5, RVC, double-float ABI
Size of this header:                  64 (bytes)
Size of program headers:              0 (bytes)
Number of program headers:            0
Size of section headers:              64 (bytes)
Number of section headers:            10
Section header string table index: 9
```

Sections

ELF Header

.shstrtab

.strtab

.symtab

.comment

.bss

.note

.data

.rela.text

.text

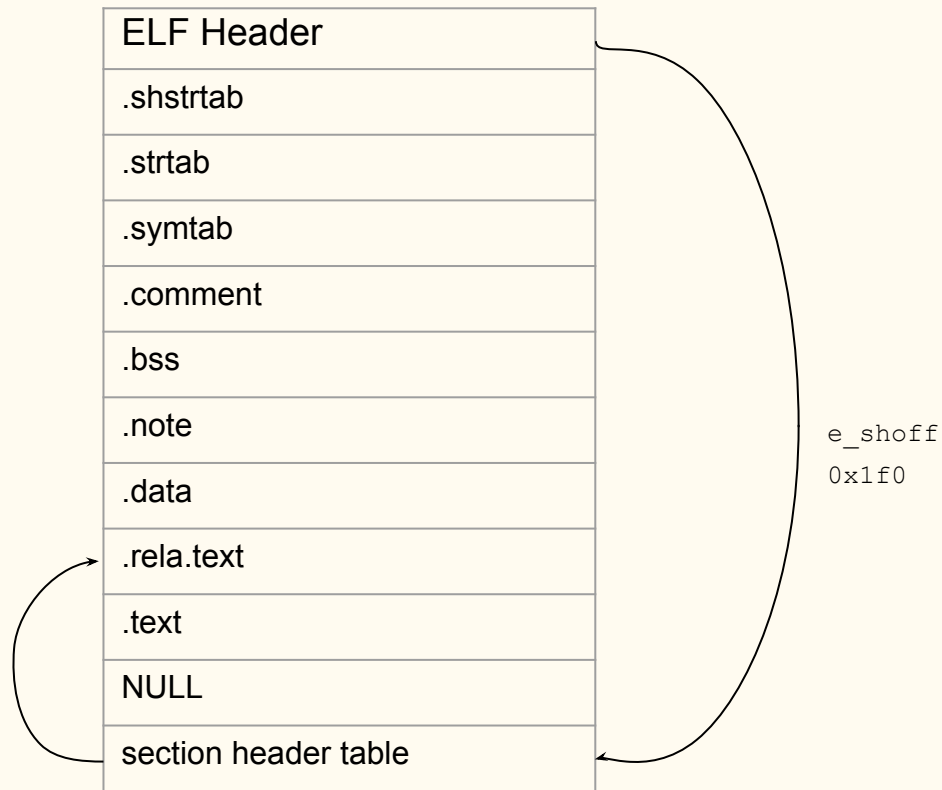
NULL

section header table

Layout of main.o

## Table look-up for `.rela.text` section

$e\_shoff - (2 * e\_shentsize)$   
 $= 0x170$



Layout of `main.o`

# Sections

- Name
- Type
- Virtual address
- File Offset
- Size
- Flag
- Alignment

```
typedef struct
{
    uint32_t  sh_name;      /* Section name (string tbl index) */
    uint32_t  sh_type;      /* Section type */
    uint64_t  sh_flags;     /* Section flags */
    uint64_t  sh_addr;      /* Section virtual addr at execution */
    uint64_t  sh_offset;    /* Section file offset */
    uint64_t  sh_size;      /* Section size in bytes */
    uint32_t  sh_link;      /* Link to another section */
    uint32_t  sh_info;      /* Additional section information */
    uint64_t  sh_addralign; /* Section alignment */
    uint64_t  sh_entsize;   /* Entry size if section holds table */
} Elf64_Shdr;
```

riscv-glibc/elf/elf.h



# Relocation section

Name	Type	Virtual address	Link	Info
------	------	-----------------	------	------

```
Section Headers:
[Nr] Name          Type          Address          Offset    Size          EntSize          Flags    Link    Info    Align
[ 0]              NULL          0000000000000000 00000000 0000000000000000 0000000000000000
[ 1] .text           PROGBITS      0000000000000000 00000040 0000000000000026 0000000000000000 AX      0      0      2
[ 2] .rela.text      RELA          0000000000000000 00000170 0000000000000030 0000000000000018 I       7      1      8
[ 3] .data           PROGBITS      0000000000000000 00000066 0000000000000000 0000000000000000 WA      0      0      1
[ 4] .bss            NOBITS        0000000000000000 00000066 0000000000000000 0000000000000000 WA      0      0      1
[ 5] .comment        PROGBITS      0000000000000000 00000066 0000000000000013 0000000000000001 MS      0      0      1
[ 6] .note.GNU-stack PROGBITS      0000000000000000 00000079 0000000000000000 0000000000000000
[ 7] .symtab          SYMTAB        0000000000000000 00000080 00000000000000d8 0000000000000018 8      7      8
[ 8] .strtab          STRTAB        0000000000000000 00000158 0000000000000013 0000000000000000 0      0      1
[ 9] .shstrtab        STRTAB        0000000000000000 000001a0 000000000000004a 0000000000000000 0      0      1
```

```
readelf -S main.o
```

# Relocation table entry

- location where a relocation will take place
- index of symbol
- relocation type
- addend

```
Relocation section '.rela.text' at offset 0x170 contains 2 entries:
  Offset          Info          Type           Sym. Value      Sym. Name + Addend
000000000000000e  00080000000012 R_RISCV_CALL    0000000000000000 add + 0
000000000000000e  00000000000033 R_RISCV_RELAX                   0
```

rela.text section in main.o

```
/* Relocation table entry with addend (in section of type SHT_RELA). */
typedef struct
{
    uint64_t r_offset;    /* Address */
    uint64_t r_info;      /* Relocation type and symbol index */
    int64_t  r_addend;    /* Addend */
} Elf64_Rela;
```

riscv-glibc/elf/elf.h

## Table look-up in `.rela.text` section

the relocation section is linked to two other sections. One of them is the symbol table, where the symbols that will be relocated are held.

`sh_link=7`

`sh_info=1`

ELF Header
.shstrtab
.strtab
.symtab
.comment
.bss
.note
.data
.rela.text
.text
NULL
section header table

# Relocation table entry

- location where a relocation will take place
- index of symbol
- relocation type
- addend

```
Relocation section '.rela.text' at offset 0x170 contains 2 entries:
  Offset          Info          Type           Sym. Value      Sym. Name + Addend
000000000000000e  00080000000012 R_RISCV_CALL    0000000000000000 add + 0
000000000000000e  00000000000033 R_RISCV_RELAX                   0
```

rela.text section in main.o

```
/* Relocation table entry with addend (in section of type SHT_RELA). */
typedef struct
{
    uint64_t r_offset;    /* Address */
    uint64_t r_info;      /* Relocation type and symbol index */
    int64_t  r_addend;    /* Addend */
} Elf64_Rela;
```

riscv-glibc/elf/elf.h

# Relocation type

Enum	ELF Reloc Type	Description	Field	Calculation	Details
0	R_RISCV_NONE	None			
1	R_RISCV_32	Runtime relocation	<i>word32</i>	$S + A$	
2	R_RISCV_64	Runtime relocation	<i>word64</i>	$S + A$	
3	R_RISCV_RELATIVE	Runtime relocation	<i>wordclass</i>	$B + A$	
4	R_RISCV_COPY	Runtime relocation			Must be in executable; not allowed in shared library
5	R_RISCV_JUMP_SLOT	Runtime relocation	<i>wordclass</i>	$S$	Handled by PLT unless <code>LD_BIND_NOW</code>
6	R_RISCV_TLS_DTPMOD32	TLS relocation	<i>word32</i>	$S \gg \text{TLSINDEX}$	

## Relocations type in RISC-V

- 0~58 Defined
- 59~255 Reserved

# Relocation type

Enum	ELF Reloc Type	Description	Field	Calculation	Details
18	R_RISCV_CALL	PC-relative call	J-Type	$S + A - P$	Macros <code>call</code> , <code>tail</code>
51	R_RISCV_RELAX	Previous reloc can be relaxed			

Relocations type *R\_RISCV\_CALL* and *R\_RISCV\_RELAX*

J type field:

- the immediate field in a J-type instruction

**R\_RISCV\_CALL** is associated with pairs of instructions (**AUIPC+JALR**) generated by the **CALL** or **TAIL** pseudoinstructions.

```
auipc x1, <hi20bits>
jalr  x0, x1, <lo12bits>
```

# Calculation symbols

Concat main.o and add.o

```
riscv64-unknown-linux-gnu-ld -relocatable main.o add.o -o main_concat.o
```

```
riscv64-unknown-linux-gnu-readelf -rs main_concat.o
```

```
Relocation section '.rela.text' at offset 0x1c8 contains 2 entries:
  Offset          Info           Type           Sym. Value      Sym. Name + Addend
000000000000000e 0008000000012 R_RISCV_CALL 000000000000026 add + 0
000000000000000e 0000000000033 R_RISCV_RELAX                0
```

Symbol table '.symtab' contains 10 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000000000	0	SECTION	LOCAL	DEFAULT	1	
2:	0000000000000000	0	SECTION	LOCAL	DEFAULT	3	
3:	0000000000000000	0	SECTION	LOCAL	DEFAULT	4	
4:	0000000000000000	0	SECTION	LOCAL	DEFAULT	5	
5:	0000000000000000	0	SECTION	LOCAL	DEFAULT	6	
6:	0000000000000000	0	FILE	LOCAL	DEFAULT	ABS	main.i
7:	0000000000000000	0	FILE	LOCAL	DEFAULT	ABS	add.c
8:	000000000000026	40	FUNC	GLOBAL	DEFAULT	1	add
9:	0000000000000000	38	FUNC	GLOBAL	DEFAULT	1	_start

## Calculation Symbols

The following table provides details on the variables used in relocation calculation:

Variable	Description
A	Addend field in the relocation entry associated with the symbol
P	Position of the relocation
S	Value of the symbol in the symbol table

$P = 0xe$

$S = 0x26$

$A = 0x0$

$S+A-P = 0x18$

# Symbols calculation for R\_RISCV\_CALL

```
main_concat.o:      file format elf64-littleriscv

Disassembly of section .text:

0000000000000000 <_start>:
 0: 1141      addi    sp,sp,-16
 2: e422      sd      s0,8(sp)
 4: 0800      addi    s0,sp,16
 6: 0ff00513  li      a0,255
 a: 06400593  li      a1,100
 e: 00000097  auipc   ra,0x0
                   e: R_RISCV_CALL add
                   e: R_RISCV_RELAX *ABS*
12: 000080e7  jalr    ra # e <_start+0xe>
16: 05d00893  li      a7,93
1a: 00000073  ecall
1e: 0001      nop
20: 6422      ld      s0,8(sp)
22: 0141      addi    sp,sp,16
24: 8082      ret

0000000000000026 <add>:
26: 1101      addi    sp,sp,-32
28: ec22      sd      s0,24(sp)
2a: 1000      addi    s0,sp,32
```

P = 0xe

S = 0x26

A = 0x0

S+A-P = 0x18

riscv64-unknown-linux-gnu-objdump -Dr main\_concat.o



# R\_RISCV\_RELEX

## Linker optimization -- Relocation Relaxation

```
main_concat.o:      file format elf64-littleriscv

Disassembly of section .text:

0000000000000000 <_start>:
 0: 1141      addi    sp,sp,-16
 2: e422      sd      s0,8(sp)
 4: 0800      addi    s0,sp,16
 6: 0ff00513  li      a0,255
 a: 06400593  li      a1,100
 e: 00000097  auipc   ra,0x0
                   e: R_RISCV_CALL add
                   e: R_RISCV_RELAX *ABS*
12: 000080e7  jalr    ra,#e<_start+0xe>
16: 05d00893  li      a7,93
1a: 00000073  ecall
1e: 0001      nop
20: 6422      ld      s0,8(sp)
22: 0141      addi    sp,sp,16
24: 8082      ret

0000000000000026 <add>:
26: 1101      addi    sp,sp,-32
28: ec22      sd      s0,24(sp)
2a: 1000      addi    s0,sp,32
```

riscv64-unknown-linux-gnu-objdump -Dr main\_concat.o

```
a.out:      file format elf64-littleriscv

Disassembly of section .text:

000000000000100b0 <_start>:
100b0: 1141      addi    sp,sp,-16
100b2: e422      sd      s0,8(sp)
100b4: 0800      addi    s0,sp,16
100b6: 09b00513  li      a0,155
100ba: 06400593  li      a1,100
100be: 014000ef  jal     ra,100d2 <add>
100c2: 05d00893  li      a7,93
100c6: 00000073  ecall
100ca: 0001      nop
100cc: 6422      ld      s0,8(sp)
100ce: 0141      addi    sp,sp,16
100d0: 8082      ret

000000000000100d2 <add>:
```

riscv64-unknown-linux-gnu-objdump -Dr a.out

## Without R\_RISCV\_RELEX

Compile without linker optimization with flag `-mno-relax`

```
riscv64-unknown-linux-gnu-gcc -c -mno-relax main.c -o main.o
```

```
riscv64-unknown-linux-gnu-gcc -c add.c -o add.o
```

```
riscv64-unknown-linux-gnu-ld main.o add.o
```

```
a.out:      file format elf64-littleriscv

Disassembly of section .text:

0000000000100b0 <start>:
100b0:      1141          addi    sp,sp,-16
100b2:      e422        sd      s0,8(sp)
100b4:      0800          addi    s0,s0,16
100b6:      09b00513     li      a0,155
100ba:      06400593     li      a1,100
100be:      00000097     auipc   ra,0x0
100c2:      018080e7     jalr    24(ra) # 100d6 <add>
100c6:      05d00893     li      a7,93
100ca:      00000073     ecall
100ce:      0001        nop
100d0:      6422        ld      s0,8(sp)
100d2:      0141          addi    sp,sp,16
100d4:      8082        ret

0000000000100d6 <add>:
```

0x18

```
riscv64-unknown-linux-gnu-objdump -Dr a.out
```

What's the address of entry point?

```
lnan95@lnan95-P65-67HSH:~/leetcode/test_linker$ riscv64-unknown-linux-gnu-ld -verbose | grep _start
ENTRY(_start)
  PROVIDE ( _executable_start = SEGMENT_START("text-segment", 0x10000)); . = SEGMENT_START("text-segment", 0x10000) + SIZEOF_HEADERS;
    PROVIDE_HIDDEN ( __rela_iplt_start = .);
    PROVIDE_HIDDEN ( __idata_start = .);
  PROVIDE_HIDDEN ( __preinit_array_start = .);
  PROVIDE_HIDDEN ( __init_array_start = .);
  PROVIDE_HIDDEN ( __fini_array_start = .);
  __bss_start = .;
```

Program entry address

# ABI (Application Binary Interface)

- High-level ABI (Itanium C++ ABI)
  - Name demangling
  - Data Layout
- Generic ABI (UNIX System V ABI)
  - System call
  - Function Calling Sequence
- Processor-specific ABI (RISC-V psABI)
  - Register Convention
  - Procedure Calling Convention
  - Code models
  - C type details

```
_start:  
    addi    sp,sp,-16  
    sd      s0,8(sp)  
    addi    s0,sp,16  
    li      a0,255  
    li      a1,100  
    call    add  
    li      a7,93  
    ecall  
    nop
```

main.s

Name demangling in C++ ABI

call \_Z3addii

Riscv64-unknown-linux-gnu-c++filt \_Z3addii

>> add(int, int)

# Register Convention

## Integer Register Convention

Name	ABI Mnemonic	Meaning	Preserved across calls?
x0	zero	Zero	-- (Immutable)
x1	ra	Return address	No
x2	sp	Stack pointer	Yes
x3	gp	Global pointer	-- (Unallocatable)
x4	tp	Thread pointer	-- (Unallocatable)
x5-x7	t0-t2	Temporary registers	No
x8-x9	s0-s1	Callee-saved registers	Yes
x10-x17	a0-a7	Argument registers	No
x18-x27	s2-s11	Callee-saved registers	Yes
x28-x31	t3-t6	Temporary registers	No

## Floating-point Register Convention

Name	ABI Mnemonic	Meaning	Preserved across calls?
f0-f7	ft0-ft7	Temporary registers	No
f8-f9	fs0-fs1	Callee-saved registers	Yes*
f10-f17	fa0-fa7	Argument registers	No
f18-f27	fs2-fs11	Callee-saved registers	Yes*
f28-f31	ft8-ft11	Temporary registers	No

RISC-V psABI: Integer Register Convention

frame pointer: s0

integer registers tp and gp should not be modified

# Named ABIs

This specification defines the following named ABIs:

- **ILP32**: Integer calling-convention only, hardware floating-point calling convention is not used (i.e. ELFCLASS32 and EF\_RISCV\_FLOAT\_ABI\_SOFT).
- **ILP32F**: ILP32 with hardware floating-point calling convention for FLEN=32 (i.e. ELFCLASS32 and EF\_RISCV\_FLOAT\_ABI\_SINGLE).
- **ILP32D**: ILP32 with hardware floating-point calling convention for FLEN=64 (i.e. ELFCLASS32 and EF\_RISCV\_FLOAT\_ABI\_DOUBLE).
- **ILP32E**: **ILP32E calling-convention** only, hardware floating-point calling convention is not used (i.e. ELFCLASS32, EF\_RISCV\_FLOAT\_ABI\_SOFT, and EF\_RISCV\_RVE).
- **LP64**: Integer calling-convention only, hardware floating-point calling convention is not used (i.e. ELFCLASS64 and EF\_RISCV\_FLOAT\_ABI\_SOFT).
- **LP64F**: LP64 with hardware floating-point calling convention for FLEN=32 (i.e. ELFCLASS64 and EF\_RISCV\_FLOAT\_ABI\_SINGLE).
- **LP64D**: LP64 with hardware floating-point calling convention for FLEN=64 (i.e. ELFCLASS64 and EF\_RISCV\_FLOAT\_ABI\_DOUBLE).
- **LP64Q**: LP64 with hardware floating-point calling convention for FLEN=128 (i.e. ELFCLASS64 and EF\_RISCV\_FLOAT\_ABI\_QUAD).

Default  
ABI

RISC-V ELF psABI specification named ABI

# C type details

- **LP64, LP64F, LP64D, and LP64Q**: use the following type sizes and alignments (based on the LP64 convention):

Type	Size (Bytes)	Alignment (Bytes)
bool/_Bool	1	1
char	1	1
short	2	2
int	4	4
long	8	8
long long	8	8
__int128	16	16
void *	8	8
float	4	4
double	8	8
long double	16	16
float _Complex	8	4
double _Complex	16	8
long double _Complex	32	16

C type sizes and alignments in RISC-V 64

# Procedure Calling Convention

- Integer Calling Convention
- Hardware Floating-point Calling Convention



# Integer Calling Convention

The base integer calling convention provides eight argument registers, a0-a7, the first two of which are also used to return values.

size < XLEN

passed in a single argument register

Aggregate  
/Scalars

```
li    a0, imm  
call func
```

XLEN=32bits for ILP32\* ABI

XLEN=64bits for ILP64\* ABI

\*Assume registers are available

# Integer Calling Convention

The base integer calling convention provides eight argument registers, a0-a7, the first two of which are also used to return values.

Aggregate /Scalars	size $\leq$ XLEN	passed in a single argument register
	size in (XLEN, 2*XLEN]	passed in a pair of argument registers for low bits and high bits

```
ld  a0, %lo
ld  a1, %hi
call func
```

XLEN=32bits for ILP32\* ABI

XLEN=64bits for ILP64\* ABI

\*Assume registers are available

# Integer Calling Convention

The base integer calling convention provides eight argument registers, a0-a7, the first two of which are also used to return values.

Aggregate /Scalars	size $\leq$ XLEN	passed in a single argument register
	size in (XLEN, 2*XLEN]	passed in a pair of argument registers for low bits and high bits
	size $>$ 2*XLEN	passed by reference and are replaced in the argument list with the address

XLEN=32bits for ILP32\* ABI

XLEN=64bits for ILP64\* ABI

\*Assume registers are available

# Hardware Floating-point Calling Convention

The hardware floating-point calling convention adds eight floating-point argument registers, fa0-fa7, the first two of which are also used to return values.

size  $\leq$  FLEN

passed in a single argument floating point register

Scalars

Otherwise

passed according to the integer calling convention

FLEN is the width of a floating-point register in the ABI

ILP64 and ILP32 ABI do not support hardware floating-point calling convention

\*Assume registers are available

# Hardware Floating-point Calling Convention

The hardware floating-point calling convention adds eight floating-point argument registers, fa0-fa7, the first two of which are also used to return values.

Aggregate with mixed variables	One floating point	passed as though it were a standalone floating-point real
	Two floating point	passed in two floating-point registers
	One floating-point and one integer	passed in a floating-point register and an integer register
	Otherwise	passed according to the integer calling convention

\*Assume registers are available, and  
sizes of fp are all less or equal then FLEN

# Code models

The code model determines what these addressing templates look like, and thus which relocations are emitted.

- Small
- Medium(default)
- Compact

Recall:

```
main_concat.o:      file format elf64-littleriscv

Disassembly of section .text:

0000000000000000 <_start>:
 0: 1141          addi    sp,sp,-16
 2: e422          sd      s0,8(sp)
 4: 0800          addi    s0,sp,16
 6: 0ff00513      li      a0,255
 a: 06400593      li      a1,100
 e: 00000097      auipc   ra,0x0
                   e: R_RISCV_CALL add
                   e: R_RISCV_RELAX *ABS*
12: 000080e7      jalr    ra # e <_start+0xe>
16: 05d00893      li      a7,93
1a: 00000073      ecall
1e: 0001          nop
20: 6422          ld      s0,8(sp)
22: 0141          addi    sp,sp,16
24: 8082          ret

0000000000000026 <add>:
26: 1101          addi    sp,sp,-32
28: ec22          sd      s0,24(sp)
2a: 1000          addi    s0,sp,32
```

riscv64-unknown-linux-gnu-objdump -Dr main\_concat.o

# Code models

- Small
  - address the whole RV32 address space or the lower 2 GiB of the RV64 address space
  - `lui` and `ld` or `st` instruction pair
- Medium
- Compact

Relocation Type:

R\_RISCV\_HI20

R\_RISCV\_LO12\_I

```
Relocation section '.rela.text' at offset 0x178 contains 4 entries:
  Offset      Info          Type           Sym. Value  Sym. Name + Addend
000000000006  00070000001a R_RISCV_HI20  0000000000000000 global_symbol + 0
000000000006  000000000033 R_RISCV_RELAX 0000000000000000 0
00000000000a  00070000001b R_RISCV_LO12_I 0000000000000000 global_symbol + 0
00000000000a  000000000033 R_RISCV_RELAX 0000000000000000 0
```

```
long global_symbol[2];
```

```
int main() {
    return global_symbol[0] != 0;
}
```

```
// -mcmmodel=medlow
```

```
lui    a5,%hi(global_symbol)
```

```
ld     a0,%lo(global_symbol)(a5)
```

```
snez   a0,a0
```

# Code models

- Small
  - address the whole RV32 address space or the lower 2 GiB of the RV64 address space
  - `lui` and `ld` or `st` instruction pair
- Medium
  - address the range between -2 GiB and +2 GiB from its position
  - `auipc` and `ld` or `st`
- Compact

Relocation Type:

R\_RISCV\_PCREL\_HI20

R\_RISCV\_PCREL\_LO12\_I

```
Relocation section '.rela.text' at offset 0x198 contains 4 entries:
Offset      Info          Type           Sym. Value  Sym. Name + Addend
000000000006 000000000017 R_RISCV_PCREL_HI20 0000000000000000 global_symbol + 0
000000000006 000000000033 R_RISCV_RELAX      0
00000000000a 000600000018 R_RISCV_PCREL_LO12_I 0000000000000006 .L0 + 0
00000000000a 000000000033 R_RISCV_RELAX      0
```

```
long global_symbol[2];
```

```
int main() {
    return global_symbol[0] != 0;
}
```

```
// -mcmmodel=medany
.LA0: auipc a5, %pcrel_hi(global_symbol)
addi a5,a5,%pcrel_lo(.LA0)
ld a5,0(a5)
snez a5,a5
```



# Code models

- Small
  - address the whole RV32 address space or the lower 2 GiB of the RV64 address space
  - `lui` and `ld` or `st` instruction pair
- Medium
  - address the range between -2 GiB and +2 GiB from its position
  - `auipc` and `ld` or `st`
- Compact
  - address the whole 64-bit address space
  - `lui` and `addi`

# Reference

1. The RISC-V Instruction Set Manual Volume I: User-Level ISA, page 67
2. RISC-V ELF psABI specification
3. SYSTEM V APPLICATION BINARY INTERFACE MIPS RISC Processor, page 4-16
4. CppCon 2019: Louis Dionne “The C++ ABI From the Ground Up”
5. CppCon 2017: Michael Spencer “My Little Object File: How Linkers Implement C++”
6. Itanium C++ ABI
7. All Aboard, Part 4: The RISC-V Code Models